

Loop Fusion

Andre Meyering, Niklas Korz, David Hector

Jan 28 2020

C++ Practice

Agenda

Development goals

Implementation

Testing

Benchmarks

Usage & Installation

Used Features

Outlook

Development goals

Development goals

```
for (size_t i = 0; i < a.size(); ++i) {  
    a[i] = b[i] + c[i];  
}  
for (size_t i = 0; i < a.size(); ++i) {  
    d[i] = a[i] + f[i];  
}
```



```
for (size_t i = 0; i < a.size(); ++i) {  
    a[i] = b[i] + c[i];  
    d[i] = a[i] + f[i];  
}
```

- Possible Cache inefficiency
- Hard to read

- Variables in cache may be reused
⇒ Possibly faster execution

Implementation

How do we fuse loops? - Overview

1. User provides invocables (e.g. lambda functions) and loop boundaries to loopers
2. **looper_union** object is created which contains all functions
3. All overlapping parts of the loop boundaries are put in individual **looper** objects
4. Loopers are executed one after the other, keeping the function sequence

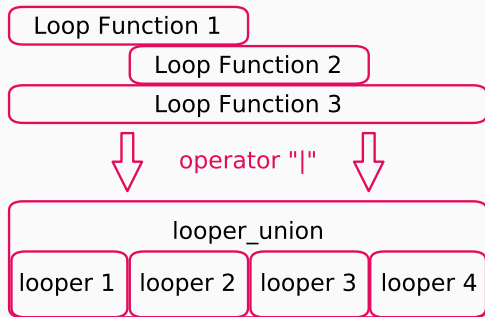


Figure 1: Schematic Loop Fusion

Detailed Implementation - Usage

```
std::array<std::vector<size_t>, 4> vec;  
const auto fill_0 = [&vec](size_t) { vec[0].push_back(0); };  
const auto fill_1 = [&vec](size_t) { vec[1].push_back(1); };  
const auto fill_2 = [&vec](size_t) { vec[2].push_back(2); };  
const auto fill_3 = [&vec](size_t) { vec[3].push_back(3); };  
  
auto l1 = loop_to(100) | fill_0;  
auto l2 = loop_from_to(10, 90, fill_1) | fill_2;  
auto l3 = loop(range{0, 100}) | fill_3  
auto merged = l1 | l2 | l3;  
merged.run();
```

Detailed Implementation - The `looper` class

General Idea

- `looper` object contains functions with the **same** loop boundaries
- Functions are stored in `std::tuple`
- Tuple is “expanded” in the for loop

Detailed Implementation - The **looper** class (simplified code)

```
template <typename Iterator, typename... F>
class looper {
public:
    constexpr looper(range<Iterator> _range, std::tuple<F...> _functions) noexcept
        : bounds { _range }
        , functions(_functions) {};

    constexpr void run()
    {
        run_loops(std::index_sequence_for<F...> {}, bounds.start, bounds.end);
    }
    // [...] | operators for merging
    // [...] sanity checks, types, etc.
public:
    common::basic_range<Iterator> bounds;
    std::tuple<F...> functions;
private:
    template <std::size_t... Idx>
    constexpr void run_loops(std::index_sequence<Idx...>, Iterator start, Iterator end)
    {
        for (Iterator i = start; i != end; ++i) {
            (std::get<Idx>(functions)(i), ...);
        }
    }
};
```

Parameter pack stored as std::tuple

Common loop boundaries

Fold expression

3 Implementations

3 Implementations

- Compile-Time-Only Merging
- Runtime Merging
- Main Range Merging

Common

- `looper_union` object contains individual ranges and function indexes
- Appending a function with `|` creates a new type

looper

- Range is template parameter
- Functions are template parameters

looper_union

- Merging in | operator
- Class contains functions and ranges with indexes for the used functions in the tuple

looper

- Runtime **range** passed as constructor argument
- **run()** method similar to compile time **looper**

looper_union

- Merging in **|** operator
- **looper** objects are stored in vector containing **std::variant**
- Running individual **looper** with **std::visit**

looper_union

- Works with arbitrary ranges computed at runtime
- Runtime overhead due to iteration over `std::vector`
- But: only when switching between ranges, functions per looper are still inlined
- `std::variant` becomes more complex with each added looper function
 - ⇒ $2^N - 1$ types in the `looper` variant
 - ⇒ All popular C++ compilers implement `std::visit` with constant time
- May result in larger binaries and more expensive calls to `std::visit`

Main Range Implementation

looper

- Runtime **range** passed as constructor argument
- **run()** method similar to compile time **looper**

looper_union

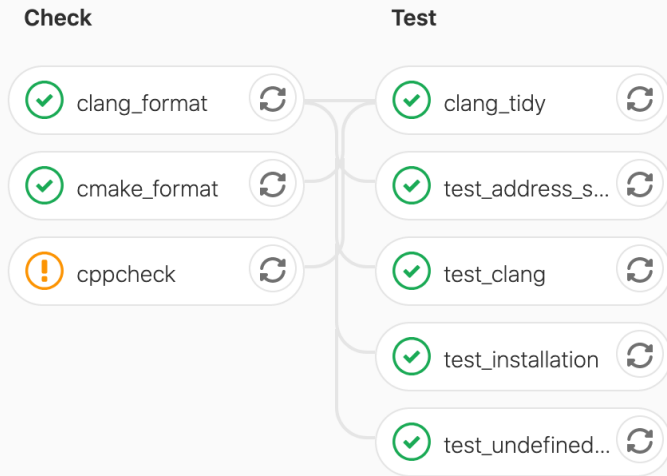
- Merging in **run()** method
- Only largest common range (intersection) is merged



Testing

We check every Merge Request against multiple tools using GitLab CI:

- Compile using GCC + Clang
- cppcheck
- Address & Undefined Behaviour Sanitizer
- clang-tidy, clang-format
- Unit Tests using Catch2



LCOV - code coverage report

Current view: [top level](#)





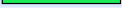
Test: [loop_fusion Test Coverage](#)

Date: [2020-01-27 18:21:29](#)

Lines:

Functions:

Hit	Total	Coverage
401	413	97.1 %
334	418	79.9 %

Directory	Line Coverage ↕			Functions ↕	
/opt/looper /tests		95.3 %	244 / 256	100.0 %	60 / 60
common		100.0 %	6 / 6	100.0 %	4 / 4
compiletime		100.0 %	36 / 36	100.0 %	61 / 61
runtime		100.0 %	60 / 60	57.5 %	100 / 174
simple		100.0 %	55 / 55	91.6 %	109 / 119

Generated by: [LCOV version 1.14-6-g40580cd](#)

Note: Harder than anticipated

GCC 9.2 does not work with lcov 1.13: Had to use latest master from GitHub!

Down the rabbit hole...

- We even found a compiler bug in GCC 9.2 (which is fixed in trunk)
- Our CI warned us before committing to master
- Clang and MSVC didn't crash!
- Demo: <https://godbolt.org/z/Q64owR>

```
during GIMPLE pass: fre
<source>: In lambda function:
<source>:307:1: internal compiler error: Segmentation fault
  307 | }
      | ^
Please submit a full bug report,
with preprocessed source if appropriate.
See <https://gcc.gnu.org/bugs/> for instructions.
Compiler returned: 1
```

Benchmarks

Benchmarks for same range loop fusion

- Three different loops
- Three equal ranges
- Implementations use very similar looper class for equal range fusion
- Comparison: Handwritten loops vs execution through looper class

```
for (size_t i = 0; i < 1000000; ++i) {  
    c[i] = a[i] + b[i];  
}
```

```
for (size_t i = 0; i < 1000000; ++i) {  
    d[i] = c[i] * a[i];  
}
```

```
for (size_t i = 0; i < 1000000; ++i) {  
    c[i] = d[i] - a[i];  
}
```

Benchmarks for same range loop fusion




Figure 2: GCC 9.2 (mp-skl2s24c)

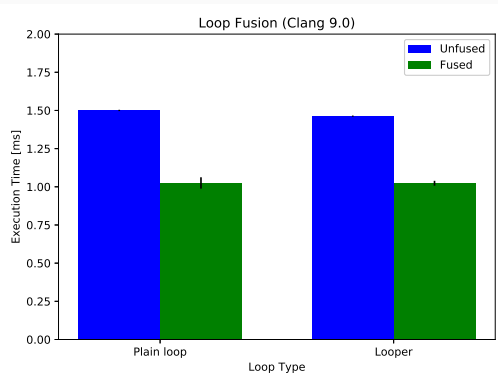


Figure 3: Clang 9.0 (mp-skl2s24c)

Benchmarks for loop unions of different ranges

- Four different loops, four partially overlapping ranges
- Seven ranges after complete merging
- Six ranges after main-range merging (only intersection of all four ranges is added)
- Comparison: Handwritten loops vs execution through united loop classes

```
for (size_t i = 0; i < 1000000; ++i) {  
    d[i] = b[i];  
    c[i] = a[i] + b[i];  
}  
for (size_t i = 5000; i < 400000; ++i) {  
    d[i] = c[i] * a[i];  
}  
for (size_t i = 50000; i < 600000; ++i) {  
    c[i] = d[i] - a[i];  
}  
for (size_t i = 300000; i < 900000; ++i) {  
    d[i] += a[i];  
}
```

Benchmarks for loop unions of different ranges




Figure 4: GCC 9.2 (mp-skl2s24c)

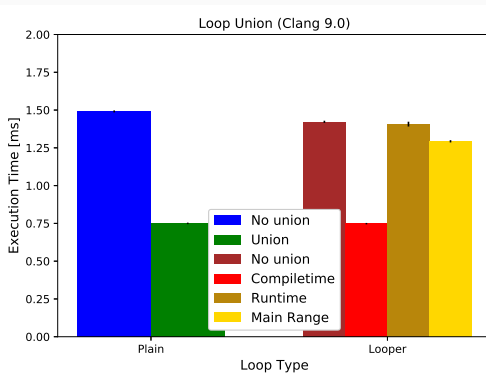


Figure 5: Clang 9.0 (mp-skl2s24c)

Note: Clang does a better job at vectorization and loop-unrolling.

Note: The assembly code for hand-fused and compile-time-fused is the same!

Usage & Installation

- We use CMake as our build system generator
 - `install` target using GNU standard directories
 - `loop_fusionConfig.cmake`, etc. are generated
 - `find_package(loop_fusion)` is enough
- Godbolt example using our library¹:



¹<https://explorer.ameyering.de/z/hc43Tz>

```
cmake --build . --target install
```

```
|— include
|   |— loop_fusion
|       |— basic_looper.hpp
|       |— ...
|— share
|   |— loop_fusion
|       |— cmake
|           |— loop_fusionConfig.cmake
|           |— loop_fusionConfigVersion.cmake
|           |— loop_fusionTargets.cmake
```

Used Features

Used Features

- Variadic Templates
- Parameter Packs + Fold Expressions
- Partial Template Specialization
- `std::variant` for runtime loop fusion
- `constexpr` for Compile time computation

Outlook

Using Iterators

- Not just integral types
- Also iterators!
- Would work now but issues for merging

```
using iterator = typename std::vector<int>::iterator;  
std::vector<int> a(100, 0);  
auto l = [&](iterator i) { *i = 10; };  
loop(basic_range<iterator>{a.begin(), a.end()}, l).run();
```

More arguments than just the iterator variable

- May be useful for Functors
- Workaround with Lambdas + Captures

```
auto l = [](ArgT1, ArgT2, size_t) { /* [...] */ };  
auto arguments = std::tuple<ArgT1, ArgT1>{ /* [...] */ };  
loop(range{0,10}, arguments, func1, func1).run();
```


Thank you for your attention

Backup Slides

C++20

- `std::vector` set to become `constexpr` in C++20²

²<http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2019/p1004r2.pdf>

RandomRangeGenerator

A `RandomRangeGenerator` with the helper function `random_range(start, end)` allowed easier testing of the mergers.

RandomRangeGenerator

```
const auto range_1 = GENERATE(  
    take(50, random_range(0, 1000))  
);
```

loop_fusion 0.1.0

loop_fusion

Main Page

Related Pages

Classes ▾

Files ▾

 Search

▼ loop_fusion

▼ Loop Fusion

► Goal

Compiler Support

Third Party Libraries

CMake Modules

► Benchmarks

► Code Quality Tools

► Code Conventions

► Design Decisions

Doxygen Documentation

► How to merge loops

Documentation

Terminology

Usage

► Classes

Loop Fusion

11 Loop fusion I a.k.a loop_fusion is a project for the course "C++ Practice WS19" from Andre Meyering, Niklas Korz and David Hector.

Goal

We try to create a header-only library that allows users to manually fuse/merge loops.

What is loop fusion?

Assume you have the following code:

```
for (size_t i = 0; i < a.size(); ++i) {  
    a[i] = b[i] + c[i];  
}  
// [...] code that is independent from the loop  
for (size_t i = 0; i < a.size(); ++i) {
```